


Limited Preemptive Scheduling



Preemptive scheduling


Most of work on scheduling has been focused on fully preemptive systems, because they allow higher responsiveness:



Disadvantages of preemptions

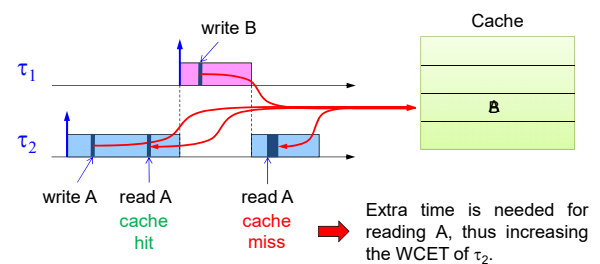
However, each preemption has a cost:

- **Context switch cost:** time taken by the scheduler to suspend the running task, switch the context, and dispatch the new incoming task.




Cache related preemption delay

CRPD: delay introduced by high priority tasks that evict cache lines containing data used in the future:




WCET

Task executing alone (or non preemptively) on a single CPU:



Task experiencing preemptions by higher priority tasks:




$$WCET_i = \begin{cases} C_i^{NP} & \text{Task executing alone} \\ C_i = C_i^{NP} + CRPD & \text{Task experiencing preemptions} \end{cases}$$



Preemption indirect costs

- **Pipeline cost:** time to flush the pipeline when a task is interrupted and to refill it when task is resumed.
- **Bus cost:** time spent waiting for the bus due to additional conflicts with I/O devices, caused by extra accesses to the RAM for the extra cache misses.



Preemption indirect costs

➤ Additional preemptions: the extra execution time also increases the number of preemptions:

Preemption cost

➤ WCETs may increase up to 35% in the presence of preemptions (less efficiency):

+35%

➤ WCETs become also more variable (less predictability):

Influence on WCETs

- As a consequence, WCETs estimations for preemptive tasks are
 - higher
 - less predictable (highly variable)

Advantages of NP scheduling

- It reduces context-switch overhead:
 - making WCETs smaller and more predictable.
- It simplifies the access to shared resources:
 - No semaphores are needed for critical sections
- It reduces stack size:
 - Task can share the same stack, since no more than one task can be in execution
- It allows achieving zero I/O Jitter:
 - $\text{finishing_time} - \text{start_time} = C_i$ (constant)

Advantages of NP scheduling

In fixed priority systems can improve schedulability:

$$U = \frac{2}{5} + \frac{4}{7} \cong 0.97$$

RM

NP-RM

Disadvantages of NP scheduling

- In general, NP scheduling reduces schedulability introducing blocking delays in high priority tasks:

Disadvantages of NP scheduling

- The utilization bound under non preemptive scheduling drops to zero:

$$U = \frac{\epsilon}{T_1} + \frac{C_2}{\infty} \rightarrow 0$$

Non preemptive anomalies

double speed

deadline miss

Non-preemptive analysis

Analysis of non-preemptive systems is more complex, because the largest response time may not occur in the first job after the critical instant.

Self-pushing phenomenon

High priority jobs activated during non-preemptive execution of lower priority tasks are pushed ahead and introduce higher delays in subsequent jobs of the same task.

Non-preemptive analysis

Hence, the analysis of τ_i must be carried out for multiple jobs, until all tasks with priority $\geq P_i$ are completed.

NOTE

Analysis can reduce to the first job of each task if and only if

- the task set is feasible under preemptive scheduling;
- All deadlines are less than or equal to periods.

Response time analysis

(for preemptively feasible task sets with $D \leq T$)

$$WO_i = \max_k \{s_{ik} - r_{ik}\}$$

$$B_i = \max_{P_j < P_i} \{C_j\}$$

$hp(i)$ set of tasks with priority higher than P_i
 $lp(i)$ set of tasks with priority lower than P_i

WO_i = Worst-case Occupied time: due to blocking B_i from $lp(i)$ tasks and interference I_i from $hp(i)$ tasks.

Response time analysis

(for preemptively feasible task sets with $D \leq T$)

$$WO_i = \max_k \{s_{ik} - r_{ik}\}$$

NOTE: the end of I_i cannot coincide with the activation of a higher priority task, because it would increase I_i .

Hence $\lceil x \rceil + 1$ must be used instead of $\lceil x \rceil$

$$WO_i = B_i + \sum_{k=1}^{i-1} \left(\left\lceil \frac{WO_i}{T_k} \right\rceil + 1 \right) C_k$$

$$R_i = WO_i + C_i$$

Response time analysis

(for preemptively feasible task sets with $D \leq T$)

$$\begin{cases} WO_i^{(0)} = B_i + \sum_{k=1}^{i-1} C_k \\ WO_i^{(s)} = B_i + \sum_{k=1}^{i-1} \left(\left\lfloor \frac{WO_i^{(s-1)}}{T_k} \right\rfloor + 1 \right) C_k \end{cases}$$

Stop when $WO_i^{(s)} = WO_i^{(s-1)}$

$$R_i = WO_i + C_i$$

19

Trade-off solutions

The following solutions can be adopted to balance between the two extreme approaches:

- **Preemption Thresholds**
 - Allow preemption only to tasks with high importance
- **Deferred Preemptions**
 - Allow preemption only after a given time interval
- **Fixed Preemption Points**
 - Allow preemption only at given points in the task code

Preemption Thresholds (PT)

Each task has two priorities:

- P_i **nominal priority**: used to enqueue the task in the ready queue and to preempt
- θ_i **threshold priority**: used for task execution ($\theta_i \geq P_i$)

A task τ_i can be preempted by τ_h only if $P_h > \theta_i$

Unfeasible task set

But feasible with PT

NOTE:
The same feasible schedule is obtained by splitting τ_3 in two non preemptive chunks: $q_{31} = 2, q_{32} = 3$

Response time analysis (PT)

$$s_{ik} = B_i + (k-1)C_i + \sum_{h: P_h > P_i} \left(\left\lfloor \frac{s_{ik}}{T_h} \right\rfloor + 1 \right) C_h$$

$$f_{ik} = s_{ik} + C_i + \sum_{h: P_h > \theta_i} \left(\left\lceil \frac{f_{ik}}{T_h} \right\rceil - \left\lfloor \frac{s_{ik}}{T_h} \right\rfloor - 1 \right) C_h$$

Deferred Preemption

Each task can defer preemption up to q_i

$$B_i = \max_{P_j < P_i} \{q_j\}$$

Interesting problem

Given a preemptively feasible task set, find the **longest non-preemptive interval** Q_i for each task that still preserves schedulability.

- Under EDF → Baruah - ECRTS 2005
- Under Fix. Pr. → Gang-Buttazzo, RTCSA 2009

Often, high priority tasks have $Q_i = C_i$, meaning that they can execute fully non preemptively.

Blocking tolerance

To compute Q_i , we need to find the maximum blocking time that can be tolerated by a task, called **blocking tolerance** (β_i):

$$\left(\sum_{k=1}^i \frac{C_k}{T_k} \right) + \frac{B_i}{T_i} \leq U_{\text{lub}}(i)$$

$$\beta_i = T_i \left[U_{\text{lub}}(i) - \sum_{k=1}^i U_k \right]$$

A simple bound for Q_i

The **longest non preemptive interval** Q_i is related with the maximum blocking time β_i that can be tolerated by higher priority tasks.

It must be $\forall i \quad B_i \leq \beta_i$

where $B_i = \max_{P_j < P_i} \{q_j\}$

Hence: $\forall i \quad \max_{P_j < P_i} \{q_j\} \leq \beta_i$

A simple bound for Q_i

$$\forall i \quad \max_{P_j < P_i} \{q_j\} \leq \beta_i$$

$$\begin{cases} i=1 & \max \{q_2, q_3, q_4\} \leq \beta_1 \\ i=2 & \max \{q_3, q_4\} \leq \beta_2 \\ i=3 & q_4 \leq \beta_3 \end{cases}$$

$$\begin{cases} i=1 & q_2 \leq \beta_1 \\ i=2 & q_3 \leq \min\{\beta_1, \beta_2\} \\ i=3 & q_4 \leq \min\{\beta_1, \beta_2, \beta_3\} \end{cases}$$


A simple bound for Q_i

$$\begin{aligned} i=1 & \quad Q_1 = \infty \\ i \geq 2 & \quad Q_i = \min\{Q_{i-1}, \beta_{i-1}\} \end{aligned}$$

Using Q_i

Once Q_i is computed, it can be used as follows:

- Partition each task into a set of NP regions no larger than Q_i inserting suitable preemption points.
- Incapsulate critical sections into NP regions, avoiding complex concurrency control protocols.



Fixed Preemption Points (FPP)

- Each task τ_i is divided in m_i chunks: $q_{i,1} \dots q_{i,m_i}$
- It can only be preempted between chunks

$$B_i = \max_{p_j < p_i} \{q_j^{\max}\}$$

Example

Let: τ_1 be fully non preemptive: $q_{11} = C_1 = 3$
 τ_2 consisting of 2 NP chunks: $q_{21} = 1, q_{22} = 3, C_2 = 4$
 τ_3 be fully non preemptive: $q_{31} = C_3 = 2$

Note that:

- The worst case response time of τ_2 does not occur in the first instance.
- The interference on τ_2 is larger than $B_2 + C_1$.

Response Time Analysis (FPP)

- Must be carry out up to the busy period of each task.

Level-i busy period

It is the interval in which the processor is busy executing tasks with priority higher than or equal to P_i , including blocking times.

Response Time Analysis (FPP)

Level-i busy period

It can be computed as the shortest interval that satisfies:

$$L_i = B_i + \sum_{h: P_h \geq P_i} \left\lceil \frac{L_i}{T_h} \right\rceil C_h$$

up to job N_i :

$$N_i = \left\lceil \frac{L_i}{T_i} \right\rceil$$

Initial value can be: $L_i^{(0)} = B_i + \sum_{h: P_h \geq P_i} C_h$

Response Time Analysis (FPP)

$$s_{ik} = B_i + (k-1)C_i + (C_i - q_i^{last}) + \sum_{h: P_h > P_i} \left(\left\lfloor \frac{s_{ik}}{T_h} \right\rfloor + 1 \right) C_h$$

$$f_{ik} = s_{ik} + q_i^{last}$$

$$R_{ik} = f_{ik} - (k-1)T_i$$

$$R_i = \max_{k \in [1, N_i]} \{R_{ik}\}$$

NOTE:
$$\begin{cases} s_{ik}^{(0)} = (k-1)T_i + (C_i - q_i^{last}) \\ N_i = \left\lfloor \frac{L_i}{T_i} \right\rfloor \end{cases}$$

Response Time Analysis (FPP)

```

for ( $i=1$  to  $n$ ) {
     $N_i = \left\lfloor L_i / T_i \right\rfloor$ 
     $s_{ik}^{(0)} = (k-1)T_i + (C_i - q_i^{last})$ 
     $k = 1$ 
    do {
         $s_{ik} = B_i + (k-1)C_i + (C_i - q_i^{last}) + \sum_{h: P_h > P_i} \left( \left\lfloor \frac{s_{ik}}{T_h} \right\rfloor + 1 \right) C_h$ 
         $f_{ik} = s_{ik} + q_i^{last}$ 
         $R_{ik} = f_{ik} - (k-1)T_i$ 
        if ( $R_{ik} > R_i$ ) then  $R_i = R_{ik}$ 
        if ( $R_i > D_i$ ) then return(UNFEASIBLE)
         $k++$ 
    } while ( $k \leq N_i$ )
}
return(FEASIBLE)

```

Special cases

- **Fully non preemptive scheduling**

$$\begin{cases} q_i^{last} = C_i \\ B_i = \max_{P_j < P_i} \{C_j\} \end{cases}$$
- **Deferred Preemption**

$$\begin{cases} q_i^{last} = 0 \\ B_i = \max_{P_j < P_i} \{Q_j\} \end{cases}$$

Final remarks

- **Preemption Thresholds** are easy to specify, but it is difficult to predict the **number of preemptions** and **where** they occur \Rightarrow large preemption overhead
- **Deferred Preemption** allows bounding the **number of preemptions** but it is difficult to predict **where** they occur. Note that the analysis assumes $q_i^{last} = 0$
- **Fixed Preemption Points** allow more control on preemptions and can be selected on purpose (e.g., to minimize overhead, stack size, and reduce WCETs).
 - A large final chunk in τ_i reduces the interference from hp-tasks (hence R_i), but creates more blocking to hp-tasks.